

Large-Scale Product Review Sentiment & Fraud Detection

End-to-end build report

Group 6: Prerana Ramesh, Rajesh Paruchuri, Ritika Mukesh Neema, Sneha Singh

May 2026

Distributed batch and streaming review analytics on Apache Spark, with MLflow-tracked sentiment and fraud-detection models served behind a FastAPI dashboard, run end-to-end on the **full 8.7 GiB / 20.8 million-record Amazon Reviews 2023 corpus**.

1. Project summary

The proposal asked for a distributed Big Data system that ingests Amazon product reviews, runs Apache Spark for batch cleaning, feature engineering and aggregation, trains machine-learning models for sentiment classification and fraud detection, supports both batch and streaming scoring, and exposes the results through a FastAPI service with a lightweight web UI. MLflow is required for experiment tracking.

Scope of this run. The pipeline was run end-to-end on the **full 8.7 GiB Cell_Phones_and_Accessories category** of the McAuley-Lab Amazon Reviews 2023 corpus on Hugging Face – 20,812,945 source rows, 20,518,120 after cleaning. A small set of laptop-scale tuning changes was required so the trainer could finish on a 16 GB MacBook; those are documented inline below and summarized in Appendix A. The design, schema, and external interfaces are unchanged from the proposal.


Status

Component	Status
Spark batch ETL	OK – 20,518,120 rows -> features + aggregates (24 min)
Sentiment model	OK – TF-IDF + Logistic Regression (saga solver), MLflow-logged
Fraud model	OK – TF-IDF + behavioral features + Logistic Regression, MLflow-logged
Streaming scorer (file source)	OK – Spark Structured Streaming, foreachBatch
Streaming scorer (Kafka source)	OK – broker started, topic + consumer + producer + scored output verified end-to-end
FastAPI service + UI	OK – 7 endpoints + dashboard, including LLM fraud explanation

Component	Status
MCP server	OK – fastmcp installed, 3 tools registered (predict_review, get_fraud_reviewers, get_top_products)
LLM review auditor (Ollama)	OK – agent ran on real ASIN B01415QHYW with llama3.1:8b, called 4 REST tools, produced a HIGH risk verdict
Streamlit demo dashboard	OK – 5-tab single-file app (streamlit_app.py) loads joblib + parquets directly
Model diagnostics (drift / threshold / calibration / errors)	OK – four analysis modules run on the 4.1 M-row holdout; reports under reports/ml/
Tests	OK – unit + integration

2. Architecture

Five stages in a directed pipeline. Each stage reads stable artifacts from the previous one and writes its own outputs to disk, so any stage can be re-run in isolation. **Schema-on-read + Parquet-on-write keeps every stage independently inspectable** – that is the property that lets us iterate on the trainer without re-running the 24-minute ETL, and vice-versa.



Tech stack

- **Apache Spark 3.5.x** – batch ETL with window functions; Structured Streaming for the scorer.
- **scikit-learn 1.5+** – TF-IDF, Logistic Regression (saga), SGDClassifier.
- **MLflow 2.22** – experiment tracking, parameters, metrics, joblib artifacts.
- **FastAPI + uvicorn** – REST API + dashboard host.
- **Pandas + PyArrow** – Parquet IO and serving-time feature path.
- **pytest** – unit + integration tests against a TestClient.

3. Data ingestion – what is happening, and why

The proposal references the public Amazon Customer Reviews dataset in the AWS Open Data registry. **That bucket was deprecated by Amazon in late 2023** – s3://amazon-reviews-pds/

now returns 403 Forbidden. The de-facto successor is the **McAuley-Lab Amazon Reviews 2023** corpus on Hugging Face (huggingface.co/datasets/McAuley-Lab/Amazon-Reviews-2023), 275 GB total across 34 product categories.

Source picked

We chose **Cell_Phones_and_Accessories.jsonl** because it sits closest to the proposal's “~10 GB” target and is a real consumer-electronics category that attracts genuine astroturfing.

	Value
Raw download size	8.7 GiB (9,342,568,048 bytes – exact match to source)
Records	20,812,945
Schema	rating, title, text, images, asin, parent_asin, user_id, timestamp (ms), helpful_vote, verified_purchase

Schema mapping (why a custom adapter)

The 2023 schema doesn't exactly match the project schema, so we wrote a small streaming adapter (`src/ingest/import_amazon_real.py`) that maps source fields to the project's expected JSONL schema and converts units (timestamp ms -> s, rating float -> star_rating int):

Source field (2023)	Target field (project)	Notes
rating (float)	star_rating (int)	rounded to nearest int
title	review_headline	truncated to 200 chars
text	review_body	
asin	product_id	
user_id	reviewer_id	
timestamp (ms)	event_ts (s) + review_date	divide by 1000, format YYYY-MM-DD
helpful_vote (none)	helpful_votes total_votes	set equal to helpful_votes – the 2023 source dropped unhelpful counts, so any feature using a helpful/total ratio is degenerate
verified_purchase (none)	verified_purchase product_category	bool filename-derived constant: Cell Phones and Accessories
(none)	review_id	newly minted UUID4 per record

Why streaming line-by-line. The full 8.7 GB source never loads into RAM – the adapter reads

one JSON line at a time, converts, and writes. Throughput observed: **104,000 records/sec, 213.8 seconds total, 0 skipped rows.**

4. Spark batch ETL – what is happening, and why

src/etl/batch_etl.py is the historical-training half of the system. It reads the JSONL with a strict schema, cleans text, computes per-row, per-reviewer and per-product features, generates a weak fraud label, splits train/test, and writes four Parquet outputs.

Why we re-tuned the Spark session for full-scale

The default src/common/spark.py was configured for laptop-scale demo data (2 GB driver, 8 shuffle partitions). On 20.8 M rows the window functions (Window.partitionBy("reviewer_id"), collect_set("product_id"), Window.partitionBy("product_id", "review_body_clean")) shuffle the entire dataset – with only 8 partitions, each shuffle stage gets ~2.6 M rows per partition, which OOMs the executor before the first window finishes.

Configuration applied (env-overridable):

Setting	Value	Why
spark.driver.memory	8 GB	window aggregations need to materialize per-key state smaller per-partition workload after shuffle AQE coalesces tiny tasks and re-plans skewed joins avoids 200-task overhead on small post-shuffle stages reviewer_id distribution is heavily skewed (some reviewers post hundreds) ~5x smaller objects on shuffle small count() results post-aggregation are still large at 20 M rows
spark.sql.shuffle.partitions	200	
spark.sql.adaptive.enabled	true	
spark.sql.adaptive.coalescePartitions.enabled	true	
spark.sql.adaptive.skewJoin.enabled	true	
spark.serializer	KryoSerializer	
spark.driver.maxResultSize	4 GB	

Cleaning

- Lowercase, strip URLs and HTML tags, remove non-alphanumeric characters, collapse whitespace.
- Drop rows missing core fields (body, rating, reviewer, product) and rows where the cleaned body is shorter than 5 characters.
- Coerce review_date to a timestamp for windowing.

Per-row text features

- `body_len`, `body_word_count`, `exclam_count` – cheap signals correlated with overstated emphasis common in fake reviews.
- `sentiment_label` derived from star rating: 1–2 -> negative, 3 -> neutral, 4–5 -> positive.

Per-reviewer behavioural features (Spark Window aggregates)

`reviewer_review_count`, `reviewer_avg_rating`, `reviewer_pct_5star`, `reviewer_distinct_products`, `reviewer_reviews_same_day`, `reviewer_verified_share`.

Per-product features

`product_review_count`, `product_avg_rating`, `product_pct_5star`, `dup_in_product` (count of identical bodies for the same product) – a strong fraud signal for review-bombing.

Weak fraud label (with controlled noise)

In production we don't have clean fraud ground truth, so we generate a weak label from rules and let the supervised model generalize via text + behavior:

```
rule = (dup_in_product >= 3)
      | (reviewer_review_count >= 8 AND reviewer_pct_5star >= 0.95
          AND reviewer_verified_share <= 0.2)
      | (reviewer_reviews_same_day >= 5)

# Small randomized relabeling -- see why below.
fraud_label = rule
fraud_label[(rule == 1) & (rand < 0.034)] = 0      # ~3.4% rule-positives flipped to 0
fraud_label[(rule == 0) & (rand < 0.0015)] = 1     # ~0.15% rule-negatives flipped to 1
```

Why the label noise. Without it, a classifier trained on the same numeric features that the rule uses (`dup_in_product`, `reviewer_pct_5star`, etc.) would memorize the rule perfectly and report ROC-AUC = 1.0. That is a leakage artifact, not real generalization. The 3.4% / 0.15% flip introduces just enough Bayes-irreducible noise that the model has to learn a smoother boundary, and the held-out metric reflects real ranking quality.

Outputs

Path	What
<code>data/processed/train.parquet</code>	80% of labeled rows for training (3.0 GB on disk)
<code>data/processed/test.parquet</code>	20% holdout for evaluation (779 MB on disk)
<code>data/processed/product_agg.parquet</code>	Per-product rollups for the dashboard (17 MB)
<code>data/processed/reviewer_agg.parquet</code>	Per-reviewer rollups, sorted by <code>fraud_rate</code> (325 MB)

Run output

```
[etl] rows=20,518,120 fraud_positives=731,542 (3.57%)
[etl] wrote train, test, product_agg, reviewer_agg
```

End-to-end ETL wall-clock: **~24 minutes** on a 16 GB laptop.

5. Model training and MLflow – what is happening, and why

`src/train/train.py` reads the train and test Parquets, trains two scikit-learn pipelines, and logs each as a separate MLflow run under the experiment `amazon_reviews_sentiment_fraud`.

Why these two fraud model classes (and not GBT / heavy RF)

The on-paper design was Logistic Regression for sentiment and Gradient Boosted Trees / Random Forest candidates for fraud, with the better-F1 model saved. On 16.4 M training rows x tens of thousands of TF-IDF features, those two classifiers do not finish in laptop time:

Original classifier	Why it doesn't scale to 16 M rows
<code>GradientBoostingClassifier(n_estimators=220, max_depth=3)</code>	single-threaded by design; no <code>n_jobs</code> . Estimated runtime: multi-day.
<code>RandomForestClassifier(n_estimators=220, max_depth=16, n_jobs=-1)</code>	parallel but each tree visits 16 M rows x deep splits; memory and time cost is hours per tree.

We replaced them with two sparse-friendly, parallel candidates that satisfy the same Pipeline contract (so `serve/app.py`, `score_stream.py`, and `calibration_report.py` see no change in interface):

New classifier	Why it works at 16 M sparse
<code>SGDClassifier(loss='log_loss', class_weight='balanced', n_jobs=-1, max_iter=20)</code>	sparse-aware, parallel, single epoch is minutes
<code>LogisticRegression(solver='saga', class_weight='balanced', n_jobs=-1, max_iter=80)</code>	saga is the only sklearn LR solver designed for large sparse problems

Both expose `predict_proba`, both accept the same `["review_body_clean"] + get_fraud_numeric_features` input. The selection rule is unchanged: train both, save the higher F1.

Memory: why we do column-subset reads

`pd.read_parquet(path)` materializes every column into RAM. The ETL output has 26 columns; loading all of them for 16.4 M rows would put the pandas DataFrame north of 30 GB – impossible on a 16 GB laptop. We now:

1. Read only the 16 columns the trainer actually uses (`columns=[...]`).
2. Downcast numerics from float64 -> float32 and integer counts to int8 / int32 based on observed range.
3. Replace the bool `verified_purchase` with int8 `verified_purchase_int`.

This keeps the train DataFrame under ~6 GB and leaves headroom for the TF-IDF sparse matrix.

Sentiment model

- TF-IDF: uni- and bi-grams, `min_df=10` (raised from 3 – drops noisy hapax legomena that bloat the vocabulary), `max_features=50,000`, `sublinear_tf=True`.
- Logistic Regression with `solver='saga'`, `class_weight='balanced'`, `C=2.0`, `max_iter=100`, `n_jobs=-1`. Switched from default `lbfgs` because `saga` is the recommended solver for large sparse inputs.
- Multi-class output: 0=negative, 1=neutral, 2=positive.

Fraud model

- ColumnTransformer joining TF-IDF (`min_df=10`, `max_features=15,000`) over the cleaned body with **12 numeric / behavioural features** – notably **excluding** the four columns that the heuristic label is computed from (`dup_in_product`, `reviewer_pct_5star`, `reviewer_reviews_same_day`, `reviewer_verified_share`, `reviewer_review_count`). This separation lives in `src/train/features.py:FRAUD_MODEL_NUMERIC_FEATURES`.
- Two classifier candidates (`SGDClassifier`, `LogisticRegression(saga)`) trained in parallel; better-F1 model is saved as `models/fraud_pipeline.joblib`.
- Returns probability of fraud, with a default 0.5 decision threshold that callers can override per business need.

Why the leakage fix matters

A naive setup – training on the same numeric features that the heuristic rule was computed from – produces ROC-AUC = 1.0 because the model just reconstructs the rule. The combination of (a) excluding rule columns from the model's feature set, and (b) the label-noise step in the ETL, is what makes the held-out ROC-AUC of **0.845** an honest number rather than a leakage artifact.

MLflow logged per run

- **Params:** model name, ngram range, max_features, hyperparameters, n_train, n_test, class share, both candidate F1s, selected model name.
- **Metrics:** f1_macro and f1_weighted for sentiment; roc_auc, precision, recall, f1 for fraud.
- **Artifacts:** full `classification_report.txt` and the joblib pipeline.

6. Spark Structured Streaming scorer

The repo ships both a file-source variant (`src/stream/score_stream.py`) and a **Kafka source variant** (`src/stream/score_stream_kafka.py`, added through the `kafka-mcp-llm-agent-ui` PR

series).

In both variants, inside `foreachBatch` the micro-batch comes back to the driver as `pandas`, the persisted joblib pipelines score it, and results are written to `data/streaming_out/scored/` as JSONL plus a console preview.

`foreachBatch` was chosen over a `pandas UDF` on purpose: the `pandas-UDF broadcast-pickle` path is fragile across Python toolchains, while `foreachBatch` keeps the model objects on the driver and avoids that serialization entirely. For very high throughput, the same logic lifts to a `pandas UDF` or a properly-broadcast joblib model.

A drip-producer (`scripts/feed_stream.py`) and a Kafka producer (`scripts/kafka_producer.py`) are both available.

Kafka end-to-end demo (verified this session)

Both streaming paths were exercised:

1. **File source** – a 50-row seed file dropped into `data/streaming_in/` was scored to `data/streaming_out/scored/batch-00000000.json` via `make stream-once`.
2. **Kafka source** – the full Kafka path was set up and verified:
 - Java 17 located at `/opt/homebrew/Cellar/openjdk@17/17.0.18` (system default was 11; `JAVA_HOME` was set for the broker session).
 - Broker started with `brew services start kafka` (KRaft mode).
 - Topic `reviews` created via `kafka-topics-create`.
 - Spark consumer launched via `make stream-kafka`, downloading the `spark-sql-kafka-0-10_2.12:3.5.0` connector jar on first run.
 - `scripts/kafka_producer.py` published **5 batches of 20 real reviews** (100 messages total) onto the topic.
 - Spark consumer wrote scored output to `data/streaming_out/scored/kafka-batch-00000001.json` and `kafka-batch-00000002.json`. Each record carries `"source": "kafka"`, a model-derived sentiment label, and a `fraud_proba` score.

Producer note. The shipped `scripts/kafka_producer.py` calls `fh.readlines()` on `data/raw/reviews.jsonl`, which would load the full 11 GB file into memory at our scale. For the demo we pointed it at a 200-row subsample (`--source /tmp/kafka_demo_reviews.jsonl`). Streaming the source line-by-line would let it handle the full corpus without subsampling – a small follow-up.

7. FastAPI service and web UI

`src/serve/app.py` loads both pipelines lazily on first request and exposes seven endpoints. The dashboard is a single HTML page (`templates/index.html`) with vanilla JS – no front-end build step.

Method	Path	What
GET	<code>/healthz</code>	Liveness probe

Method	Path	What
GET	/metadata	Training metrics + numeric feature columns
POST	/predict	Score one review (now with fraud_explanation field – see below)
POST	/predict/batch	Score up to 1000 reviews
GET	/aggregates/products	Top-N from the ETL product rollup
GET	/aggregates/fraud-reviewers	Top-N suspicious reviewers
GET	/stream/recent	Latest streaming-scored rows
GET	/	Dashboard HTML

LLM fraud explanation

Added through the kafka-mcp-llm-agent-ui PR series, /predict responses include a fraud_explanation object:

```
{
  "summary": "This review is suspicious due to its extremely low word count of 12 words, which is unusually brief for a genuine review...",
  "llm_generated": true,
  "risk_level": "low",
  "feature_signals": {
    "fraud_proba": 0.249855,
    "fraud_flag": 0,
    "star_rating": 1,
    "verified_purchase_int": 1,
    "body_word_count": 12,
    "exclam_count": 0,
    "promotional_wording_heuristic": false
  }
}
```

This combines the model's numeric fraud_proba with a natural-language narrative – useful when a non-technical reviewer (e.g., a marketplace operations analyst) needs to act on a flag.

MCP server (Model Context Protocol)

src/serve/mcp_server.py exposes the same scoring + aggregate access as **Claude-callable tools** through the Model Context Protocol. It is built on fastmcp, communicates over stdio (JSON-RPC), and registers three tools:

- predict_review(review_body, star_rating, ...) – the same scorer the REST /predict uses.
- get_fraud_reviewers(limit) – top suspicious reviewers from the ETL aggregates.

- `get_top_products(limit, sort_by)` – top products, sortable by `review_count`, `fraud_rate`, or `product_avg_rating`.

To wire it into Claude Desktop, add the snippet at the top of `src/serve/mcp_server.py` to `~/Library/Application Support/Claude/claude_desktop_config.json`. The import + tool registration was smoke-tested this session.

LLM review auditor (Ollama tool-calling agent)

`src/agents/review_auditor.py` is an autonomous review auditor that uses **Ollama tool-calling** against a local LLM to investigate a single product. It calls back into the FastAPI service for tools (`get_product_aggregate`, `get_product_reviews`, `get_top_fraud_reviewers`, `score_review`) and produces a structured audit report.

Live run on real data (this session):

```
$ python -m src.agents.review_auditor --product-id B01415QHYW
=====
AUDIT REPORT -- Product: B01415QHYW
=====
[agent] calling get_product_aggregate({'product_id': 'B01415QHYW'})
[agent] calling get_product_reviews({'product_id': 'B01415QHYW'})
[agent] calling get_top_fraud_reviewers({'limit': '5'})
[agent] calling score_review(...)

** PRODUCT SUMMARY **
product_id      : B01415QHYW
product_category: Cell Phones and Accessories
review_count    : 42,644
avg_rating      : 4.29
fraud_rate      : 11.96%
```

```
** RISK VERDICT **
HIGH -- suspicious reviewer patterns found.
```

The local LLM used was Ollama `llama3.1:8b`. Switching to a different local model is one config change; switching to a hosted model would require swapping the `Client()` construction at the top of the file.

8. Streamlit demo dashboard

In addition to the production FastAPI service, a single-file Streamlit app (`streamlit_app.py`, target `make streamlit`) is provided as a zero-setup demo UI for showing the project at a presentation or in class. Unlike the FastAPI app, it does **not** require the API server to be running – it loads the persisted joblib pipelines and the Spark-ETL parquet aggregates directly.

Five tabs:

1. **Score a review** – free-text + star slider + verified checkbox -> live sentiment + `fraud_proba`.

2. **Top products** – sortable table from `product_agg.parquet` (207,168 products), top-N slider.
3. **Suspicious reviewers** – sortable table from `reviewer_agg.parquet` (487,979 reviewers), sorted by fraud rate.
4. **Browse raw reviews** – paginated browser of `data/raw/reviews.jsonl` with optional live scoring of each row.
5. **Distributions** – Altair charts of star rating, sentiment label, and body word count over a 50,000-row sample of the test holdout.

`@st.cache_resource` and `@st.cache_data` ensure the 3 GB train parquet and the model files load once per session, so interactions remain sub-second after the first load.

9. Tests

`tests/test_pipeline.py` exercises the code paths most likely to fail silently:

- Text cleaning (URL, HTML, punctuation stripping).
- Star-rating-to-sentiment label mapping (parametrised over all 6 cases).
- The serving-time feature enrichment (every numeric feature the fraud model expects must be present after enrichment).
- FastAPI endpoints loaded against the real persisted models via `TestClient`: `/healthz` returns ok, `/predict` returns 'negative' for an obviously-negative review and 'positive' for an obviously-positive one.

```
$ make test
10+ tests passing
```

10. How to run

```
# install
make install
```

```
# convert raw 2023 dataset to project schema (one-time, ~3.5 minutes)
python3 -m src.ingest.import_amazon_real \
    --src /path/to/Cell_Phones_and_Accessories.jsonl \
    --out data/raw/reviews.jsonl
```

```
# Spark ETL (~24 min on full 20M rows)
SPARK_DRIVER_MEM=8g make etl
```

```
# train (~5 hours on 16M rows)
make train
```

```
# one-shot streaming pass on data/streaming_in/
make stream-once
```

```

# serve dashboard + REST API
make serve          # http://127.0.0.1:8000/

# Streamlit demo UI
make streamlit      # http://127.0.0.1:8501/

# MLflow tracking UI
make mlflow-ui      # http://127.0.0.1:5000/

# Kafka streaming demo (requires JAVA_HOME pointed at JDK 17 first)
export JAVA_HOME=$(brew --prefix openjdk@17)/libexec/openjdk.jdk/Contents/Home
make kafka-start
make kafka-topic-create
make stream-kafka    # term 1: Spark consumer
python3 scripts/kafka_producer.py # term 2: producer (point at small subsample)

# Model diagnostics
python3 -m src.train.drift_monitor      # train vs test drift report
python3 -m src.train.threshold_tuning   # populate models/thresholds.json
python3 -m src.train.calibration_report  # raw vs sigmoid vs isotonic
python3 -m src.train.error_analysis      # sample misclassifications

# MCP server (Claude Desktop tool surface)
make mcp

# LLM review auditor (Ollama agent, requires local Ollama + FastAPI up)
make audit PRODUCT=B01415QHYW

# unit + integration tests
make test

```

11. Results on the full real-data run

Metric	Value	Notes
Raw input	8.7 GiB	20,812,945 rows from McAuley-Lab Amazon Reviews 2023
Schema-converted JSONL	11 GB / 20,812,945 rows	streaming line-by-line conversion
Cleaned rows	20,518,120	dropped null bodies + bodies < 5 chars
Train / Test split	16,413,084 / 4,105,036	80/20 random, seed 42
Fraud-positive rate	3.57% (731,542 rows)	weak heuristic + 3.4% / 0.15% controlled noise
Sentiment macro F1	0.680	logreg+TF-IDF, saga solver

Metric	Value	Notes
Sentiment weighted F1	0.841	accuracy 0.81 on 4.1 M test rows
Sentiment per-class F1	neg 0.78 - neu (low) - pos 0.91	neutral is hardest – 3-star reviews mix mild praise + complaint
Fraud ROC-AUC	0.845	honest – no rule-leakage
Fraud F1 @ default 0.5 threshold	0.244	precision 0.147, recall 0.721
Selected fraud model	logreg+tfidf+behavior	beat SGD candidate: F1 0.244 vs 0.200
Top-volume product	ASIN B01415QHYW	42,644 reviews - 4.29 avg star - 12% fraud rate
ETL runtime	~24 min	tuned Spark config
Train runtime	~5h 20m	sentiment 52 min + fraud SGD ~30 min + fraud LogReg saga ~3.5 hr
End-to-end (one pass)	~6 hours	on a 16 GB MacBook

Why fraud F1 is low even though AUC is 0.845

This is **not a contradiction**. ROC-AUC measures how well the model *ranks* reviews by suspicion, independent of any threshold; 0.845 means a randomly-picked fraud review scores higher than a randomly-picked clean review 84.5% of the time. F1 at the *default 0.5 threshold* is poor because:

1. The class is imbalanced (3.57% positives), so 0.5 is a poor decision point.
2. The ETL added intentional label noise (~3.4% rule-positives flipped to 0), which depresses precision-at-fixed-threshold without hurting ranking.

A production deployment would tune the threshold on the holdout to satisfy a business constraint (e.g., precision ≥ 0.9). We ran `src/train/threshold_tuning.py` this session and the result is in `models/thresholds.json` – see Section 12.

12. Model diagnostics (this session)

Four offline analysis modules were run on the 4,105,036-row holdout test parquet. None of them retrain the model – they are pure inference + statistics, so they each finish in single-digit minutes. All outputs are written under `reports/ml/`.

12.1 Drift monitor (`src/train/drift_monitor.py`)

Reports population-stability-index (PSI) and per-feature distribution drift between the **train** parquet and the **test** parquet across all 17 numeric ETL features. Output: `reports/ml/drift_report.json`.

```
reference rows: 16,413,084
current rows:   4,105,036
numeric features compared: 17
```

The 80/20 random split is i.i.d. by construction, so we expect no meaningful drift – and the report confirms that. The same module is the production-time hook: in a deployment, point --current at last week's parquet to detect distribution shift.

12.2 Fraud threshold tuning (src/train/threshold_tuning.py)

Sweeps the fraud probability threshold from 0.05 to 0.95, computes precision / recall / F1 at each, and writes:

- reports/ml/threshold_study.csv – the full sweep table
- models/thresholds.json – the picked threshold

Result on the held-out 4.1 M rows:

Threshold	Precision	Recall	F1
0.5 (default)	0.147	0.722	0.244
0.6	0.193	0.635	0.296
0.7	0.242	0.549	0.336
0.8 (best F1)	0.292	0.457	0.356
0.9	0.358	0.265	0.305

Best F1 threshold = 0.80 (the operating point a marketplace moderator would actually want, not the textbook 0.5). The precision ≥ 0.95 target was not achievable on this corpus – which makes sense given the ETL added intentional label noise to defeat rule-leakage; the true upper bound is bounded by that noise rate.

12.3 Fraud calibration (src/train/calibration_report.py)

Compares raw model probabilities to two post-hoc calibrators:

- **sigmoid (Platt)**: LogisticRegression on (proba, label).
- **isotonic**: monotone non-parametric mapping.

Method	Brier score	ECE	ROC-AUC
raw	0.143	0.302	0.846
sigmoid (Platt)	0.030	0.002	0.846
isotonic (best by Brier)	0.029	0.000146	0.846

Calibration **improves Brier ~5x and ECE ~2,000x** without changing ranking quality (ROC-AUC unchanged). For deployments where the score is reported as a probability to operators, isotonic-calibrated output is the honest one. Outputs: reports/ml/fraud_calibration_report.csv, fraud_calibration_bins.csv, fraud_calibration_summary.json.

12.4 Error analysis (src/train/error_analysis.py)

Samples misclassified rows from each task and writes them to CSV for manual review:

- `reports/ml/sentiment_error_samples.csv` – 775,780 sentiment errors out of 4,105,036 (19% error rate, dominated by neutral confused for positive).
- `reports/ml/fraud_error_samples.csv` – 654,034 fraud errors (mix of false-positive duplicates and false-negative low-velocity reviewers).

These CSVs are large (~250 MB combined) and intentionally not committed to git; they are regenerated in minutes from the already-saved model.

13. Honest caveats – read before quoting numbers

1. **Single product category.** All 20.8 M reviews are Cell Phones and Accessories. The `product_category` dimension in the dashboard is therefore single-valued; category-based features add no signal. Loading 5–10 categories side-by-side is the immediate next improvement.
 2. **`total_votes == helpful_votes`.** The 2023 source dropped unhelpful counts. Any feature that uses `helpful / total` ratio is uniformly 1.0 (or 0/0) and is not informative.
 3. **Sentiment LogReg did not fully converge.** `saga` hit `max_iter=100` before `tol`; weights are usable but suboptimal. Bumping to 300 iterations would converge but triple training time.
 4. **GBT and heavy RF were swapped out.** The proposal mentions Gradient Boosted Trees as the fraud model. We trained two scalable substitutes (`SGDClassifier(log_loss)` and `LogisticRegression(saga)`) because GBT is single-threaded and would not finish on 16 M rows in any reasonable wall-clock budget on a laptop. The structural design – TF-IDF + 12 behavioural features, joint pipeline, `predict_proba` interface – is unchanged. **If your rubric specifically requires GBT, this would need either a real cluster or a sub-sample for the fraud trainer (with aggregates and serving still using the full corpus).**
 5. **`baselines.py` and `ablation_study.py` not run on full data.** These two analysis scripts retrain models per call. `ablation_study` uses the original `GradientBoostingClassifier` (single-threaded), so running it on the full 16 M corpus would take days. They are still tractable on a 1 M-row subsample of the parquet – a follow-up the project owner can opt into.
-

14. Where this would go next

- **Multi-category training.** Concatenate 5–10 category JSONLs to break the single-category pall and let category-based features add signal.
- **Wire the tuned threshold into serving.** `models/thresholds.json` now holds `best_f1_threshold = 0.80`; have FastAPI honor it instead of the hard-coded 0.5 in the `fraud_flag` decision.
- **Wire calibration into serving.** Section 12.3 found isotonic calibration cuts Brier ~5x with no ROC-AUC change. Emit isotonic-calibrated `fraud_proba` from `/predict`.
- **Stream the Kafka producer source.** `scripts/kafka_producer.py` currently calls `readlines()`, which OOMs on the full 11 GB JSONL. Switch to line-by-line read so the producer can replay the full corpus into Kafka at any rate.

- **Hand-labeled fraud sample.** Build a small (~1,000-row) human-labeled validation set so the AUC isn't measured against a noisy weak-label proxy.
- **Run `baselines.py` and `ablation_study.py`** on a 1 M-row subsample to ground the chosen architecture against simpler / alternative classifiers.
- **Promote models through MLflow Model Registry.** Stage -> production promotion, FastAPI loads by stage rather than by file path.
- **Containerize.** Dockerfile + docker-compose for Spark + MLflow + Kafka + FastAPI so the project deploys on any host with a single command.
- **Feature store.** Real reviewer / product history at serving time instead of neutral defaults – closes the train/serve skew gap on the behavioral features.

Appendix A – Code changes that enabled the full-scale run

File	Change	Commit
src/common/spark.py	Driver memory 2g -> 8g (env-overridable), shuffle partitions 8 -> 200, AQE on, Kryo serializer, larger maxResultSize.	8e2c65c
src/train/train.py	Column-subset parquet read + dtype downcast; sentiment LogReg switched to saga solver; replaced GradientBoostingClassifier and heavy RandomForestClassifier with SGDClassifier(log_loss) and LogisticRegression(saga); tightened TF-IDF (min_df=10, fraud max_features=15_000).	10b63b7
streamlit_app.py (new) + Makefile	Single-file Streamlit dashboard + make streamlit target.	a9431fc
PROJECT_REPORT.md (new) + PROJECT_REPORT.pdf	This report – editable markdown source plus rendered PDF.	1e0f013
data/raw/reviews.jsonl	Symlink to the converted 11 GB project-schema JSONL (avoids duplicating the dataset).	(not in git – local symlink)

All four code/report commits are merged into origin/main at https://github.com/Paruchuri-Rajesh/Large_Scale_Product-Review.