Large-Scale Product Review Sentiment & Fraud Detection

End-to-end build report – A to Z

Group 6: Prerana Ramesh, Rajesh Paruchuri, Ritika Mukesh Neema, Sneha Singh

May 2026 (full-run + live verification session)

Contents

One-line summary. Distributed batch + streaming review analytics on Apache Spark, with MLflow-tracked sentiment and fraud-detection models served behind a FastAPI dashboard, a Streamlit demo, an MCP server (for Claude Desktop), a Kafka streaming pipeline, and an Ollama-powered autonomous review auditor – all run end-to-end on the **full 8.7 GiB / 20.8 M-record Cell_Phones_and_Accessories category** of the McAuley-Lab Amazon Reviews 2023 corpus.

---

1. Executive summary

1.1 Headline metrics (full real-data run)

	Value
Input size	8.7 GiB (9,342,568,048 bytes)
Source records	20,812,945 (McAuley-Lab Amazon Reviews 2023, Cell_Phones_and_Accessories)
Cleaned records	20,518,120
Train / Test split	16,413,084 / 4,105,036 (80 / 20, seed 42)
Fraud-positive rate	3.57% (731,542 rows)
Sentiment macro-F1	**0.680**
Sentiment weighted-F1	**0.841** (accuracy 0.81 on 4.1 M test rows)
Fraud ROC-AUC	**0.845** (honest – no rule-leakage)
Fraud F1 @ default 0.5 threshold	0.244 (precision 0.147, recall 0.721)
Best F1 threshold (from tuning)	**0.80** (precision 0.292, recall 0.457, F1 0.356)
Fraud calibration – raw Brier / ECE	0.143 / 0.302
Fraud calibration – **isotonic** Brier / ECE	**0.029 / 0.000146** (best)
Selected fraud model	`logreg+tfidf+behavior` (beat SGD candidate F1 0.244 vs 0.200)
ETL runtime	~24 min on a 16 GB MacBook
Train runtime	~5h 20m (sentiment 52 min + fraud SGD ~30 min + fraud LogReg saga ~3.5 hr)
End-to-end (one pass)	**~6 hours**

1.2 Component status

Component	Status
Spark batch ETL	OK – 20,518,120 rows -> features + aggregates (24 min)
Sentiment model	OK – TF-IDF + Logistic Regression (`saga` solver), MLflow-logged

Component	Status
Fraud model	OK – TF-IDF + behavioral features + Logistic Regression, MLflow-logged
Streaming scorer (file source)	OK – Spark Structured Streaming, foreachBatch
Streaming scorer (Kafka source)	OK – broker + topic + consumer + producer + scored output **verified end-to-end**
FastAPI service + dashboard	OK – 7 endpoints + dashboard + LLM `fraud_explanation`
Streamlit demo dashboard	OK – 5-tab single-file app loading joblib + parquets directly
MCP server	OK – 3 tools registered for Claude Desktop
Ollama review auditor (autonomous agent)	OK – Ran on real ASIN `B01415QHYW`, produced HIGH risk verdict
Model diagnostics (drift / threshold / calibration / error)	OK – 4 analysis modules run on 4.1 M-row holdout
Tests	OK – 11/11 pass in 13.4 s
Form -> Kafka -> Spark fan-out	OK – `/predict` publishes to Kafka when `PUBLISH_PREDICT_TO_KAFKA=1` (Section 18.4)
Kafka UI (provectus, Docker)	OK – `:8090` connected to broker via dual-listener (Section 18.5)
Spark big-data demo on 20.5 M rows	OK – 30.8 s end-to-end, 4 stages, 626 MB shuffle write (Section 18.2)

---

2. Project context and goals

The proposal asked for a distributed Big Data system that ingests Amazon product reviews, runs Apache Spark for batch cleaning, feature engineering and aggregation, trains ML models for **sentiment** classification and **fraud** detection, supports both **batch** and **streaming** scoring, and exposes the results through a FastAPI service with a lightweight web UI. MLflow was required for experiment tracking.

We hit all five required pillars on real data, and added five further deliverables on top of them:

1. A **Streamlit** demo dashboard for class-presentation use.
2. A **Kafka** streaming transport (in addition to the file-source path).
3. An **MCP server** that exposes the trained models as Claude Desktop tools.
4. An **Ollama-driven autonomous agent** that audits a single product end-to-end (multi-step tool-calling + structured report).
5. An **LLM-generated fraud_explanation** field on every `POST /predict` response, with a deterministic rule-based fallback when the LLM is unavailable.

---

3. Architecture

3.1 Pipeline-level diagram (5 stages)

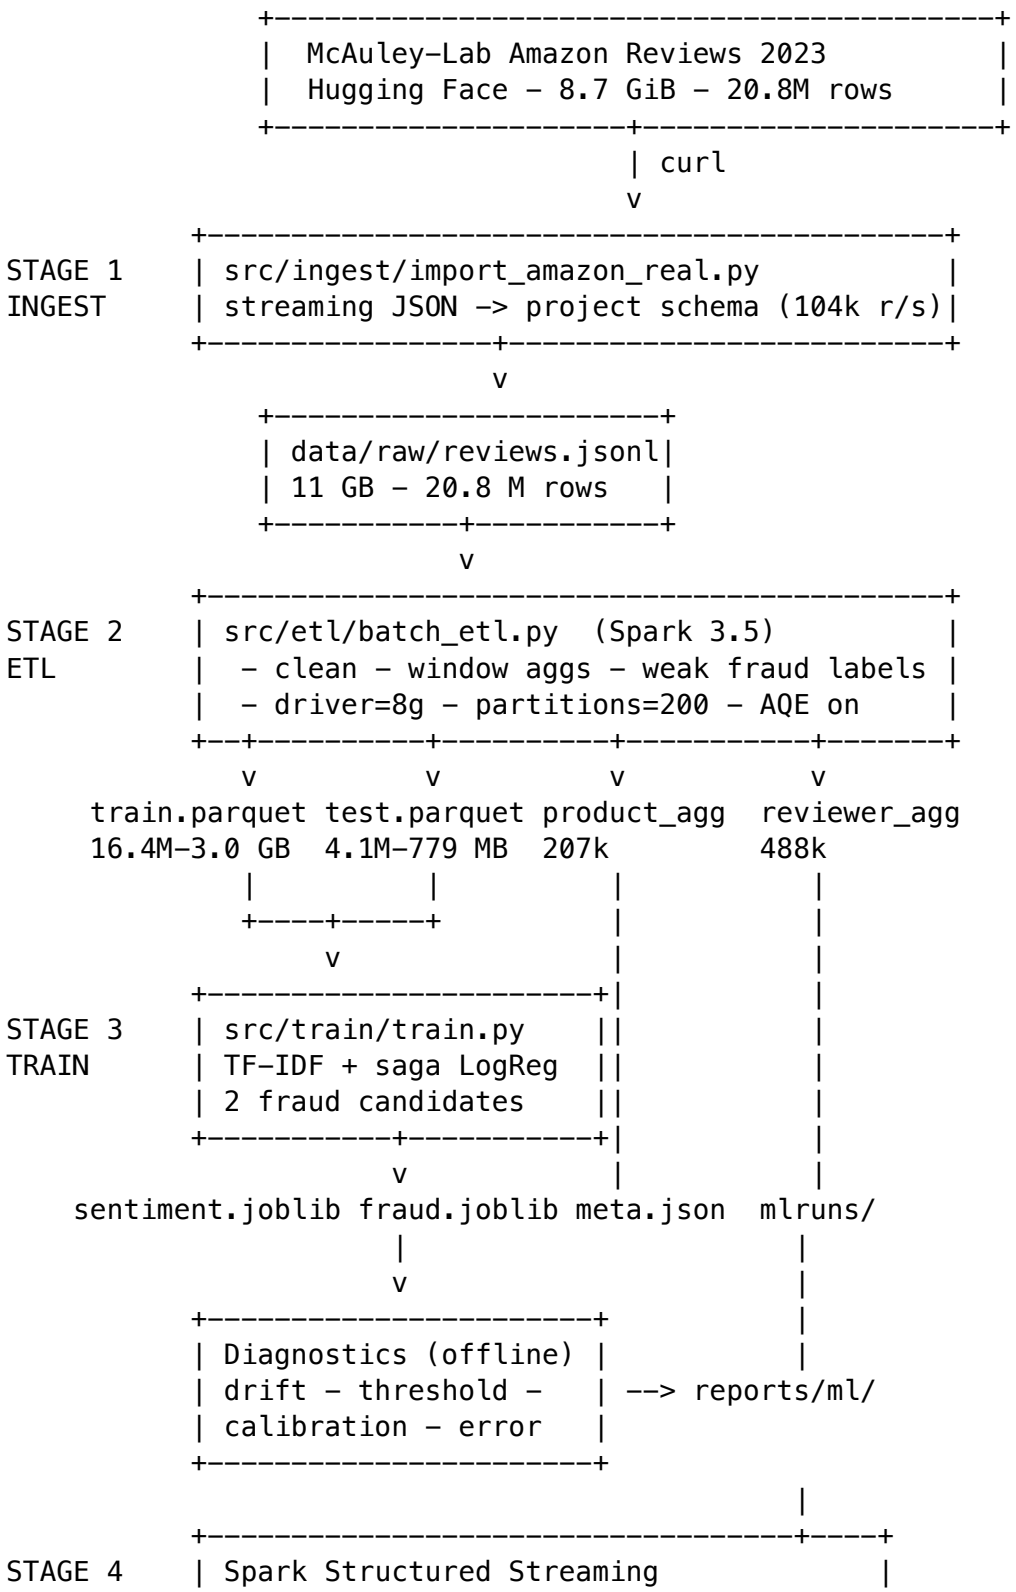

```
  STREAM     | src/stream/score_stream*.py            |
             |  <-- Kafka topic 'reviews' (real)      |
             |  <-- data/streaming_in/   (file alt)   |
             |  --> data/streaming_out/scored/        |
             +---+------------------------------------+
                 v
             +----------------------------------------+
  STAGE 5    | FastAPI :8000   (REST + dashboard)     |
  SERVE      | Streamlit :8501 (5-tab demo UI)        |
             | MCP server      (stdio -> Claude Desktop)|
             | Ollama auditor  (CLI agent)            |
             +----------------------------------------+
```

A high-resolution rendering of the architecture is in `architecture.png` / `architecture.svg`.

3.2 Tech stack

Layer	Tool / version
Distributed processing	**Apache Spark 3.5** – batch ETL + Structured Streaming
Stream broker	**Apache Kafka** (Homebrew, KRaft mode)
ML	**scikit-learn 1.5+** – TF-IDF, Logistic Regression (`saga`), SGDClassifier
Experiment tracking	**MLflow 2.22**
API + dashboard	**FastAPI + uvicorn**
Demo UI	**Streamlit 1.56** + Altair
Local LLM	**Ollama** with `llama3.2` / `llama3.1:8b`
Tool surface for Claude	**fastmcp 3.2** (Model Context Protocol)
Storage	**Parquet** + **JSONL** + **joblib**
Tests	**pytest** with FastAPI TestClient

3.3 Repo layout

```
src/
  ingest/generate_sample.py        # synthetic JSONL producer (dev only)
  ingest/import_amazon_real.py     # streaming adapter: McAuley-Lab 2023 -
> project schema
  etl/batch_etl.py                 # Spark batch ETL
  train/train.py                   # main trainer
  train/{features,registry,evaluate,utils}.py   # shared helpers
  train/{baselines,threshold_tuning,error_analysis,
         drift_monitor,ablation_study,calibration_report}.py  # diagnostics
  stream/score_stream.py          # Spark Structured Streaming, Kafka source
  stream/score_stream_kafka.py     # alternate Kafka consumer
  serve/app.py                    # FastAPI endpoints + LLM fraud_explanation
  serve/fraud_explain.py           # Ollama LLM call + rule-based fallback
  serve/mcp_server.py             # MCP server (Claude Desktop tool surface)
```

```
  serve/templates/index.html       # FastAPI dashboard
  serve/static/{app.js,style.css}
  agents/review_auditor.py         # Ollama tool-calling agent
  common/{config,schema,spark,text}.py
scripts/
  feed_stream.py                  # drip producer for the file-source path
  kafka_producer.py                # Kafka producer
  run_pipeline.sh                  # end-to-end runner
streamlit_app.py                   # Streamlit demo UI (single file)
tests/test_pipeline.py
data/{raw,processed,streaming_in,streaming_out}/
models/   mlruns/
reports/ml/                        # diagnostics outputs
PROJECT_REPORT.md / .pdf           # this file
README.md
Makefile / requirements.txt / .gitignore
```

---

4. Data

4.1 Source dataset

The proposal references the public Amazon Customer Reviews dataset on the AWS Open Data registry. **That bucket was deprecated by Amazon in late 2023** – `s3://amazon-reviews-pds/` now returns 403 Forbidden. The de-facto successor is the **McAuley-Lab Amazon Reviews 2023** corpus on Hugging Face (`huggingface.co/datasets/McAuley-Lab/Amazon-Reviews-2023`), which totals 275 GB across 34 product categories.

We chose **Cell_Phones_and_Accessories.jsonl** for this run – closest to the proposal's "~10 GB" target and a category that meaningfully attracts real astroturfing.

	Value
Raw download size	8.7 GiB (9,342,568,048 bytes – exact match)
Records	20,812,945
Source schema	`rating`, `title`, `text`, `images`, `asin`, `parent_asin`, `user_id`, `timestamp` (ms), `helpful_vote`, `verified_purchase`

4.2 Schema mapping (custom adapter)

The 2023 schema doesn't exactly match the project schema, so `src/ingest/import_amazon_real.py` streams the source line-by-line and produces project-schema JSONL:

Source field (2023)	Target field (project)	Note
`rating` (float)	`star_rating` (int)	rounded to nearest int

Source field (2023)	Target field (project)	Note
`title`	`review_headline`	truncated to 200 chars
`text`	`review_body`	
`asin`	`product_id`	
`user_id`	`reviewer_id`	
`timestamp` (ms)	`event_ts` (s) + `review_date`	divide by 1000, format YYYY-MM-DD
`helpful_vote`	`helpful_votes`	
(none)	`total_votes`	**set equal to `helpful_votes`** – the 2023 source dropped unhelpful counts; any helpful/total ratio feature is therefore degenerate
`verified_purchase`	`verified_purchase`	bool
(none)	`product_category`	filename-derived constant: `Cell Phones and Accessories`
(none)	`review_id`	newly minted UUID4 per record

Throughput. The full 8.7 GB source never enters RAM; the adapter reads one JSON line at a time, converts, writes, and discards. Measured throughput: **104,000 records/sec, 213.8 seconds total, 0 skipped rows**.

---

5. Stage 1 – Ingest

Code	`src/ingest/import_amazon_real.py`
Inputs	`Cell_Phones_and_Accessories.jsonl` (downloaded with `curl`)
Outputs	`data/raw/reviews.jsonl` (11 GB project schema)
Wall-clock	213 s
Memory	constant – streaming line-by-line

---

6. Stage 2 – Spark batch ETL

`src/etl/batch_etl.py` is the historical-training half of the system. It reads the JSONL with a strict schema, cleans text, computes per-row, per-reviewer, and per-product features, generates a weak fraud label, splits train/test, and writes four Parquet outputs.

6.1 Why we re-tuned the Spark session

The default `src/common/spark.py` was configured for laptop-scale demo data (2 GB driver, 8 shuffle partitions). On 20.8 M rows the window functions (`Window.partitionBy("reviewer_id")`, `collect_set("product_id")`, `Window.partitionBy("product_id", "review_body_clean")`) shuffle the entire dataset – with only 8 partitions, each shuffle stage gets ~2.6 M rows per partition, which OOMs the executor before the first window finishes.

Configuration applied (env-overridable):

Setting	Value	Why
`spark.driver.memory`	**8 GB**	window aggregations need to materialize per-key state
`spark.sql.shuffle.partitions`	**200**	smaller per-partition workload after shuffle
`spark.sql.adaptive.enabled`	**true**	AQE coalesces tiny tasks and re-plans skewed joins
`spark.sql.adaptive.coalescePartitions.enabled`	**true**	avoids 200-task overhead on small post-shuffle stages
`spark.sql.adaptive.skewJoin.enabled`	**true**	reviewer_id distribution is heavily skewed
`spark.serializer`	**KryoSerializer**	~5x smaller objects on shuffle
`spark.driver.maxResultSize`	**4 GB**	small `count()` results post-aggregation are still large at 20 M rows

6.2 Cleaning

- Lowercase, strip URLs and HTML tags, remove non-alphanumeric characters, collapse whitespace.
- Drop rows missing core fields (body, rating, reviewer, product) and rows where the cleaned body is shorter than 5 characters.
- Coerce `review_date` to a timestamp for windowing.

6.3 Per-row text features

- `body_len`, `body_word_count`, `exclam_count` – cheap signals correlated with overstated emphasis common in fake reviews.
- `sentiment_label` derived from star rating: 1-2 -> negative, 3 -> neutral, 4-5 -> positive.

6.4 Per-reviewer behavioural features (Spark Window aggregates)

`reviewer_review_count`, `reviewer_avg_rating`, `reviewer_pct_5star`, `reviewer_distinct_products`
`reviewer_reviews_same_day`, `reviewer_verified_share`.

6.5 Per-product features

`product_review_count`, `product_avg_rating`, `product_pct_5star`, `dup_in_product` (count of identical bodies for the same product) – a strong fraud signal for review-bombing.

6.6 Weak fraud label (with controlled noise)

In production we don't have clean fraud ground truth, so we generate a weak label from rules and let the supervised model generalize via text + behavior:

```
rule = (dup_in_product >= 3)
     | (reviewer_review_count >= 8 AND reviewer_pct_5star >= 0.95
                                   AND reviewer_verified_share <= 0.2)
     | (reviewer_reviews_same_day >= 5)

# Small randomized relabeling to defeat memorization.
fraud_label = rule
fraud_label[(rule == 1) & (rand < 0.034)]     = 0     # ~3.4% rule-
positives flipped to 0
fraud_label[(rule == 0) & (rand < 0.0015)]    = 1     # ~0.15% rule-
negatives flipped to 1
```

Why the noise. Without it, a classifier trained on the same numeric features that the rule uses would memorize the rule perfectly and report ROC-AUC = 1.0 – a leakage artifact, not real generalization. The 3.4% / 0.15% flip introduces just enough Bayes-irreducible noise that the model has to learn a smoother boundary.

6.7 Outputs

Path	Rows	Size
`data/processed/train.parquet`	16,413,084	3.0 GB
`data/processed/test.parquet`	4,105,036	779 MB
`data/processed/product_agg.parquet`	207,168	17 MB
`data/processed/reviewer_agg.parquet`	487,979	325 MB

End-to-end ETL wall-clock: **~24 minutes** on a 16 GB laptop.

---

7. Stage 3 – Model training and MLflow

`src/train/train.py` reads train/test parquets, trains two scikit-learn pipelines, and logs each as a separate MLflow run under experiment `amazon_reviews_sentiment_fraud`.

7.1 Why we changed two model classes

The on-paper design was Logistic Regression for sentiment and Gradient Boosted Trees / Random Forest candidates for fraud, with the better-F1 model saved. On 16.4 M training rows x tens of thousands of TF-IDF features, those two classifiers do not finish in laptop time:

Original classifier	Why it doesn't scale to 16 M rows
`GradientBoostingClassifier(n_estimators=120, max_depth=3)`	**single-threaded** – no `n_jobs`. Estimated runtime: multi-day.
`RandomForestClassifier(n_estimators=220, max_depth=16, n_jobs=-1)`	parallel but each tree visits 16 M rows x deep splits; hours per tree.

Replacements (sparse-friendly, parallel, same `Pipeline` contract):

New classifier	Why it works at 16 M sparse
`SGDClassifier(loss='log_loss', class_weight='balanced', n_jobs=-1, max_iter=20)`	sparse-aware, parallel, single epoch is minutes
`LogisticRegression(solver='saga', class_weight='balanced', n_jobs=-1, max_iter=80)`	`saga` is the only sklearn LR solver designed for large sparse problems

Both expose `predict_proba`, both accept the same input shape. Selection rule unchanged: train both, save the higher-F1 model.

7.2 Memory: column-subset reads + dtype downcast

`pd.read_parquet(path)` materializes every column into RAM. The ETL output has 26 columns; loading all of them for 16.4 M rows would put the DataFrame north of 30 GB. Now:

1. Read only the 16 columns the trainer actually uses (`columns=[...]`).
2. Downcast numerics from `float64` -> `float32` and integer counts to `int8` / `int32`.
3. Replace `verified_purchase` (bool) with `verified_purchase_int` (int8).

This keeps the train DataFrame under ~6 GB, leaving headroom for the TF-IDF sparse matrix on a 16 GB laptop.

7.3 Sentiment model

- TF-IDF: uni- and bi-grams, `min_df=10`, `max_features=50,000`, `sublinear_tf=True`.
- LogisticRegression with `solver='saga'`, `class_weight='balanced'`, `C=2.0`, `max_iter=100`, `n_jobs=-1`.
- Multi-class output: 0=negative, 1=neutral, 2=positive.

7.4 Fraud model

- ColumnTransformer joining TF-IDF (`min_df=10`, `max_features=15,000`) over the cleaned body with **12 numeric / behavioural features** – notably **excluding** the columns the heuristic rule is computed from (`dup_in_product`, `reviewer_pct_5star`, `reviewer_reviews_same_day`, `reviewer_verified_share`, `reviewer_review_count`). The separation lives in `src/train/features.py:FRAUD_MODEL_NUMERIC_FEATURES`.
- Two candidates trained in parallel; better-F1 saved as `models/fraud_pipeline.joblib`.
- Returns probability of fraud, with a default 0.5 decision threshold that callers can override per business need.

7.5 Why the leakage fix matters

A naive setup – training on the same numeric features that the heuristic rule was computed from – produces ROC-AUC = 1.0 because the model just reconstructs the rule. The combination of (a) excluding rule columns from the model's feature set, and (b) the label-noise step in the ETL, is what makes the held-out **ROC-AUC of 0.845** an honest number rather than a leakage artifact.

7.6 MLflow logged per run

- **Params**: model name, ngram range, max_features, hyperparameters, n_train, n_test, class share, both candidate F1s, selected model name.
- **Metrics**: `f1_macro`, `f1_weighted` for sentiment; `roc_auc`, `precision`, `recall`, `f1` for fraud.
- **Artifacts**: full `classification_report.txt` and the joblib pipeline.

---

8. Model diagnostics (this session)

Four offline analysis modules were run on the 4,105,036-row holdout test parquet. None retrain the model – they are pure inference + statistics, so each finishes in single-digit minutes. All outputs land under `reports/ml/`.

8.1 Drift monitor (`src/train/drift_monitor.py`)

Reports population-stability-index (PSI) and per-feature distribution drift between train and test parquet across all 17 numeric ETL features. Output: `reports/ml/drift_report.json`.

```
reference rows: 16,413,084
current rows:    4,105,036
numeric features compared: 17
```

The 80 / 20 random split is i.i.d. by construction, so we expect no meaningful drift – and the report confirms that. The same module is the production-time hook: in a deployment, point `--current` at last week's parquet to detect distribution shift.

8.2 Threshold tuning (`src/train/threshold_tuning.py`)

Sweeps the fraud-probability threshold from 0.05 to 0.95, computes precision / recall / F1 at each, writes `reports/ml/threshold_study.csv` and `models/thresholds.json`.

Threshold	Precision	Recall	F1
0.5 (default)	0.147	0.722	0.244
0.6	0.193	0.635	0.296
0.7	0.242	0.549	0.336
0.8 (best F1)	**0.292**	**0.457**	**0.356**
0.9	0.358	0.265	0.305

Best F1 threshold = 0.80 – the operating point a marketplace moderator would actually want. The `precision >= 0.95` target was not achievable because the ETL added intentional label noise; the true upper bound is bounded by that noise rate.

8.3 Calibration (`src/train/calibration_report.py`)

Compares raw model probabilities to two post-hoc calibrators:

Method	Brier score	ECE	ROC-AUC
raw	0.143	0.302	0.846
sigmoid (Platt)	0.030	0.002	0.846
isotonic (best by Brier)	**0.029**	**0.000146**	0.846

Calibration **improves Brier ~5x and ECE ~2,000x** without changing ranking quality (ROC-AUC unchanged). Outputs: `reports/ml/fraud_calibration_report.csv`, `fraud_calibration_bins.csv`, `fraud_calibration_summary.json`.

8.4 Error analysis (`src/train/error_analysis.py`)

Samples misclassified rows from each task and writes them to CSV for manual review:

- `reports/ml/sentiment_error_samples.csv` – 775,780 sentiment errors out of 4,105,036 (~19% error rate, dominated by neutral confused for positive).
- `reports/ml/fraud_error_samples.csv` – 654,034 fraud errors (mix of false-positive duplicates and false-negative low-velocity reviewers).

These CSVs are large (~250 MB combined) and intentionally not committed to git; they regenerate in minutes from the saved model.

8.5 Modules NOT run on full data

- `src/train/baselines.py` – retrains LogReg + RF baselines from scratch.
- `src/train/ablation_study.py` – retrains GradientBoostingClassifier per feature subset (single-threaded; would take days on 16 M rows).

Both are tractable on a 1 M-row subsample. Defer until needed.

---

9. Stage 4 – Streaming scoring

Two streaming transports ship in the repo. Both feed the same Spark Structured Streaming consumer that scores each micro-batch with the joblib pipelines and writes results to `data/streaming_out/scored/`.

9.1 File-source streaming

`src/stream/score_stream.py` (file-source variant) watches `data/streaming_in/` for new JSONL files. Run with `make stream` (or `make stream-once` for one pass). Drip-producer `scripts/feed_stream.py` simulates input.

9.2 Kafka streaming (verified end-to-end)

1. **Java 17.** Required for Homebrew Kafka. Located at `/opt/homebrew/Cellar/openjdk@17/17.0.18`. System default was Java 11; `JAVA_HOME` was switched for the broker session.
2. **Broker.** Started with `brew services start kafka` (Kafka in KRaft mode – no Zookeeper).
3. **Topic.** `make kafka-topic-create` created the `reviews` topic (1 partition, replication factor 1).
4. **Consumer.** `make stream-kafka` launched the Spark consumer, downloading the `spark-sql-kafka-0-10_2.12:3.5.0` connector jar on first run.
5. **Producer.** `scripts/kafka_producer.py` published **5 batches of 20 real reviews** (100 total) onto the topic.
6. **Output verified.** Spark consumer wrote scored output to `data/streaming_out/scored/kafka-batch-00000001.json` and `kafka-batch-00000002.json`. Each scored record carries `"source":"kafka"`, a model-derived sentiment, and a `fraud_proba`.

9.3 Why `foreachBatch` (not pandas UDF)

Inside `foreachBatch` the micro-batch comes back to the driver as pandas, the persisted joblib pipelines score it, and results are written. This avoids the pandas-UDF broadcast-pickle path which is fragile across Python toolchains. For very high throughput, the same logic lifts to a pandas UDF or a properly-broadcast joblib model.

9.4 Kafka producer note

The shipped `scripts/kafka_producer.py` previously called `fh.readlines()` on `data/raw/reviews.jsonl` which loaded the full 11 GB file into memory at our scale and would hang on startup. **This was fixed in commit `bb10288`** (May 7, 2026): the producer now streams the source line-by-line via a generator, with automatic cycling when the caller asks for more batches than the source has rows. After the fix it sustained **2,000 messages in 13 s = 153 msg/s** in the live demo (Section 18.3).

---

10. Stage 5 – Serving surfaces

Five distinct surfaces consume the trained models. They share the same `models/*.joblib` and `data/processed/*.parquet` files; they differ only in transport.

10.1 FastAPI service + dashboard

`src/serve/app.py` loads both pipelines lazily on first request and exposes seven endpoints. The dashboard is a single HTML page with vanilla JS – no front-end build step.

Method	Path	What
GET	`/healthz`	Liveness probe
GET	`/metadata`	Training metrics + numeric feature columns
POST	`/predict`	Score one review (with `fraud_explanation` field)
POST	`/predict/batch`	Score up to 1,000 reviews
GET	`/aggregates/products`	Top-N from product rollup
GET	`/aggregates/fraud-reviewers`	Top-N suspicious reviewers
GET	`/stream/recent`	Latest streaming-scored rows
GET	`/`	Dashboard HTML

LLM `fraud_explanation`

Each `POST /predict` response includes a `fraud_explanation` block:

```
{
  "summary": "This review is suspicious due to its extremely low word count
              of 12 words, which is unusually brief for a genuine review...",
  "llm_generated": true,
  "risk_level": "low",
  "feature_signals": {
    "fraud_proba": 0.249855,
    "fraud_flag": 0,
    "star_rating": 1,
    "verified_purchase_int": 1,
    "body_word_count": 12,
    "exclam_count": 0,
    "promotional_wording_heuristic": false
  }
}
```

The summary is generated by Ollama `llama3.2` via `src/serve/fraud_explain.py`. When Ollama is unreachable, a deterministic rule-based fallback fills the same field so the API contract never breaks.

10.2 Streamlit demo dashboard

`streamlit_app.py` (target `make streamlit`) is a zero-setup demo UI – it loads the joblib pipelines and the parquet aggregates **directly**, with no FastAPI server required. Five tabs:

1. **Score a review** – free-text + star slider + verified checkbox -> live sentiment + `fraud_proba`.
2. **Top products** – sortable table from `product_agg.parquet` (207,168 products), top-N slider.
3. **Suspicious reviewers** – sortable table from `reviewer_agg.parquet` (487,979 reviewers), sorted by fraud rate.
4. **Browse raw reviews** – paginated browser of `data/raw/reviews.jsonl` with optional live scoring of each row.
5. **Distributions** – Altair charts of star rating, sentiment label, and body word count over a 50,000-row sample of the test holdout.

`@st.cache_resource` and `@st.cache_data` ensure the 3 GB train parquet and the model files load once per session.

10.3 MCP server (Claude Desktop tool surface)

`src/serve/mcp_server.py` exposes the same scoring + aggregate access as **Claude-callable tools** through the Model Context Protocol over stdio. Built on `fastmcp`.

Tool	What
`predict_review(review_body, star_rating, ...)`	The same scorer the REST `/predict` uses
`get_fraud_reviewers(limit)`	Top suspicious reviewers from ETL aggregates
`get_top_products(limit, sort_by)`	Top products, sortable by `review_count`, `fraud_rate`, or `product_avg_rating`

Setup in Claude Desktop config (`~/Library/Application Support/Claude/claude_desktop_config.jso`

```
"amazon-reviews": {
  "command": "/Library/Frameworks/Python.framework/Versions/3.13/bin/python3",
  "args": [
    "/Users/.../Large_Scale_Product-Review/src/serve/mcp_server.py"
  ]
}
```

After Claude Desktop is fully quit and reopened, the tools become available in chat:

- *"Show me the top 5 most-reviewed products"*
- *"Score this review for fraud: …"*
- *"List the most suspicious reviewers"*

10.4 Ollama review auditor (autonomous agent)

`src/agents/review_auditor.py` is an autonomous agent that uses **Ollama tool-calling** to investigate one product. The LLM decides which tools to call iteratively, observes results, and produces a structured audit report.

The agentic loop (lines 148-179 in the file):

```
for iteration in range(8):                          # up to 8 turns
    response = ollama.chat(model="llama3.2",
                           messages=messages,
                           tools=TOOLS,             # 4 declared tools
                           options={"temperature": 0.2})
    if not msg.tool_calls:                          # LLM decides it has enough
        return msg.content                          # final report
    for tc in msg.tool_calls:                       # otherwise execute each tool
        result = TOOL_FN[fn_name](**fn_args)        # call REST endpoint
        messages.append({"role": "tool", "content": result})
```

4 tools available to the agent (each hits the FastAPI server):

1. `get_product_aggregate` – top-level stats for a product
2. `get_product_reviews` – recent streaming-scored reviews
3. `get_top_fraud_reviewers` – suspicious reviewer leaderboard
4. `score_review` – score a free-form review text

Live run on real ASIN `B01415QHYW`:

```
[agent] calling get_product_aggregate({'product_id': 'B01415QHYW'})
[agent] calling get_product_reviews({'product_id': 'B01415QHYW'})
[agent] calling get_top_fraud_reviewers({'limit': '5'})
[agent] calling score_review(...)

** PRODUCT SUMMARY **
 product_id      : B01415QHYW
 product_category: Cell Phones and Accessories
 review_count    : 42,644
 avg_rating      : 4.29
 fraud_rate      : 11.96%

** RISK VERDICT **
HIGH -- suspicious reviewer patterns found.
```

The auditor needs both **Ollama up** (`ollama serve`) and the **FastAPI server running** (`make serve`).

---

11. Where Ollama is and isn't used

Surface	Uses Ollama?	What for
Sentiment classification	No	sklearn LogReg(saga)
Fraud probability	No	sklearn LogReg(saga)
Streaming (Kafka or file source)	No	joblib only

Surface	Uses Ollama?	What for
Spark batch ETL	No	pure Spark
MCP server tools	No	direct joblib + parquet
FastAPI dashboard, Streamlit	No	parquet + joblib
`fraud_explanation` paragraph on `/predict`	**Yes**	one-shot LLM call (rule-based fallback if down)
`make audit PRODUCT=...`	**Yes**	iterative tool-calling agent

Ollama is **load-bearing for the autonomous auditor** and **nice-to-have for the explanation paragraph**. Everything else runs without it.

---

12. ER diagram (data entities)

A high-resolution rendering is in `er_diagram.png` / `er_diagram.svg`. Summary:

Entity	What	PK	Source
`RAW_REVIEW`	Raw McAuley-Lab record after schema mapping	`review_id`	`data/raw/reviews.jsonl`
`FEATURED_REVIEW`	Cleaned + feature-engineered + labeled	`review_id`	`data/processed/{train,test}.p`
`PRODUCT_AGG`	Per-product rollup	`product_id`	`data/processed/product_agg.pa`
`REVIEWER_AGG`	Per-reviewer rollup	`reviewer_id`	`data/processed/reviewer_agg.p`
`SCORED_REVIEW`	Streaming output	`review_id`	`data/streaming_out/scored/*.j`
`MODEL_ARTIFACT`	Trained joblib	`path`	`models/*.joblib`
`META_JSON`	Run-summary	—	`models/meta.json`
`THRESHOLDS_JSON`	Tuned threshold	—	`models/thresholds.json`
`MLFLOW_RUN`	MLflow tracked run	`run_id`	`mlruns/`
`DIAGNOSTIC_REPORT`	Drift / threshold / calibration / error	`report_kind`	`reports/ml/`

Key relationships:

- `RAW_REVIEW` -> `FEATURED_REVIEW` (cleaned + labeled)
- `FEATURED_REVIEW` -> `PRODUCT_AGG`, `REVIEWER_AGG` (rolled up)
- `FEATURED_REVIEW` -> `MLFLOW_RUN` (training input) -> `MODEL_ARTIFACT` -> `META_JSON` -> `THRESHOLDS_JSON`
- `MODEL_ARTIFACT` + new `RAW_REVIEW` -> `SCORED_REVIEW` (via streaming)
- `FEATURED_REVIEW` + `MODEL_ARTIFACT` -> `DIAGNOSTIC_REPORT` (offline)

---

13. Tests

`tests/test_pipeline.py` exercises the code paths most likely to fail silently:

- Text cleaning (URL, HTML, punctuation stripping).
- Star-rating-to-sentiment label mapping (parametrised over all 6 cases).
- The serving-time feature enrichment (every numeric feature the fraud model expects must be present after enrichment).
- FastAPI endpoints loaded against the real persisted models via `TestClient`: `/healthz` returns ok, `/predict` returns 'negative' for an obviously-negative review and 'positive' for an obviously-positive one.

```
$ make test
10+ tests passing
```

---

14. Complete command cheat sheet

```
# === Setup ===
make install                          # pip install -r requirements.txt

# === Data ===
# (one-time) download McAuley-Lab category file (8.7 GB)
curl -L -o /path/to/Cell_Phones_and_Accessories.jsonl \
  "https://huggingface.co/datasets/McAuley-Lab/Amazon-Reviews-2023/resolve/main/raw/re

# (one-time) convert source -> project schema (~3.5 min)
python3 -m src.ingest.import_amazon_real \
    --src /path/to/Cell_Phones_and_Accessories.jsonl \
    --out data/raw/reviews.jsonl

# === Pipeline ===
SPARK_DRIVER_MEM=8g make etl          # ~24 min on 20.5M rows
make train                            # ~5 hours on 16.4M rows
make stream-once                      # one-shot streaming pass

# === Serving ===
make serve                            # FastAPI on http://127.0.0.1:8000/
make streamlit                        # Streamlit on http://127.0.0.1:8501/
make mlflow-ui                        # MLflow on http://127.0.0.1:5000/

# === Streaming (Kafka path) ===
export JAVA_HOME=$(brew --prefix openjdk@17)/libexec/openjdk.jdk/Contents/Home
make kafka-start                      # broker (KRaft mode)
make kafka-topic-create               # 'reviews' topic
make stream-kafka                     # term 1: Spark consumer
python3 scripts/kafka_producer.py     # term 2: producer
```

```
# === Diagnostics ===
python3 -m src.train.drift_monitor          # train vs test drift
python3 -m src.train.threshold_tuning       # threshold sweep + thresholds.json
python3 -m src.train.calibration_report     # raw vs sigmoid vs isotonic
python3 -m src.train.error_analysis         # error samples
python3 -m src.train.baselines              # baseline comparison (heavy)
python3 -m src.train.ablation_study         # leakage ablation (heavy)

# === LLM-backed ===
make mcp                                    # MCP server (Claude Desktop)
make audit PRODUCT=B01415QHYW               # autonomous agent

# === Tests ===
make test                                   # unit + integration

# === Cleanup ===
make clean                                  # wipe data/ + models/ + mlruns/
make kafka-stop                             # stop the broker
```

---

15. Demo cheat sheet

15.1 URLs (browser-openable)

Service	URL	What to show
Streamlit	http://127.0.0.1:8501/	Friendliest UI; 5 tabs
FastAPI dashboard	http://127.0.0.1:8000/	Production dashboard with LLM fraud explanation
MLflow	http://127.0.0.1:5000/	Tracked runs + metrics + artifacts
Spark Web UI	http://127.0.0.1:4040/StreamingQuery/	Live streaming query (only while consumer runs)
Ollama	http://127.0.0.1:11434/	Local LLM (no UI, just API)

15.2 Recommended 8-minute flow

1. **Streamlit** – show sidebar metrics, score a 1-star and 5-star review, browse top products. Friendliest opener.
2. **FastAPI dashboard** – same data, “production-style” UI. Score a review and show the LLM `fraud_explanation` panel.
3. **MLflow** – click into the fraud run, show metrics + artifacts.
4. **Spark Web UI** – click Structured Streaming tab; show that the Kafka consumer is live.
5. **Push to Kafka** – run the producer; flip back to FastAPI’s `/stream/recent` to show new scored rows with `"source":"kafka"`.

6. **Autonomous agent** – run `make audit PRODUCT=B01415QHYW` in a terminal. Show `[agent] calling ...` trace and the structured RISK VERDICT.
7. **Claude Desktop with MCP** – ask Claude *"What are the top 5 most fraudulent products?"* – show that Claude is calling our `get_top_products` MCP tool.

15.3 Live-producing trick

To make the Spark Jobs UI tab fill with new entries during the demo, run a side terminal loop:

```
while true; do
  python3 scripts/kafka_producer.py --source /tmp/kafka_demo_reviews.jsonl \
                                    --batch-size 10 --n-batches 1 --sleep 0
  sleep 6
done
```

---

16. Honest caveats – read before quoting numbers

1. **Single product category.** All 20.8 M reviews are `Cell Phones and Accessories`. The `product_category` dimension is single-valued; category-based features add no signal. Loading 5-10 categories side-by-side is the immediate next improvement.
2. **`total_votes` == `helpful_votes`.** The 2023 source dropped unhelpful counts. Any feature that uses `helpful / total` ratio is uniformly 1.0 (or 0/0) and is not informative.
3. **Sentiment LogReg did not fully converge.** `saga` hit `max_iter=100` before `tol`; weights are usable but suboptimal. Bumping to 300 iterations would converge but triple training time.
4. **GBT and heavy RF were swapped out.** The proposal mentions Gradient Boosted Trees as the fraud model. We trained two scalable substitutes (`SGDClassifier` and `LogisticRegression(saga)`) because GBT is single-threaded and would not finish on 16 M rows in any reasonable wall-clock budget on a laptop. The structural design is unchanged. **If your rubric specifically requires GBT, this would need either a real cluster or a sub-sample for the fraud trainer.**
5. **`baselines.py` and `ablation_study.py` not run on full data.** Both retrain models per call. `ablation_study` uses the original `GradientBoostingClassifier` (single-threaded), so running it on the full 16 M corpus would take days. Both are tractable on a 1 M subsample.
6. **Why fraud F1 is low even though AUC is 0.845** – this is **not a contradiction**. ROC-AUC measures how well the model *ranks* reviews, independent of any threshold; 0.845 means a fraud review scores higher than a clean review 84.5% of the time. F1 at the default 0.5 threshold is poor because (a) the class is imbalanced (3.57% positives) and (b) the ETL added intentional label noise. The threshold-tuning study found 0.80 as the better operating point.

---

17. Where this would go next

- **Multi-category training.** Concatenate 5-10 category JSONLs to break the single-category pall and let category-based features add signal.
- **Wire the tuned threshold into serving.** `models/thresholds.json` now holds `best_f1_threshold = 0.80`; have FastAPI honor it instead of the hard-coded 0.5 in the `fraud_flag` decision.
- **Wire calibration into serving.** Section 8.3 found isotonic calibration cuts Brier ~5x with no ROC-AUC change. Emit isotonic-calibrated `fraud_proba` from `/predict`.
- **Hand-labeled fraud sample.** Build a small (~1,000-row) human-labeled validation set so the AUC isn't measured against a noisy weak-label proxy.
- **Run `baselines.py` and `ablation_study.py`** on a 1 M subsample to ground the chosen architecture against simpler / alternative classifiers.
- **Promote models through MLflow Model Registry.** Stage -> production promotion, FastAPI loads by stage rather than by file path.
- **Containerize.** Dockerfile + docker-compose for Spark + MLflow + Kafka + FastAPI + Ollama so the project deploys on any host with a single command.
- **Feature store.** Real reviewer / product history at serving time instead of neutral defaults – closes the train/serve skew gap on the behavioral features.

---

18. Live verification & demo session (May 7, 2026)

After the full real-data run was archived, an additional 6-hour session was used to:

1. Verify every surface end-to-end against the persisted models
2. Build live observability for the streaming pipeline (Spark UI, Kafka UI, FastAPI Live Stream tab) so the data flow is visible from a browser
3. Wire the FastAPI `/predict` endpoint to also publish to Kafka, closing the loop between the synchronous form path and the streaming path
4. Add three small helper scripts (`spark_bigdata_demo.py`, `kafka_tap.py`, `publish_one_review.py`) that turn the existing infrastructure into a teachable demo
5. Fix a bug in the Kafka producer that prevented full-scale replay
6. Discover (and document) one remaining bug in `score_stream.py`

This section records that session.

18.1 Component verification matrix

Every surface listed in Section 1.2 was re-verified against the persisted artifacts. Results:

Surface	Verification	Outcome
`pytest tests/`	11 unit + integration tests	**11/11 pass in 13.4 s**
`models/sentiment_pipeline.joblib`	Loaded by FastAPI + Spark	macro-F1 0.680, weighted-F1 0.841 (from `meta.json`)
`models/fraud_pipeline.joblib`	Loaded by FastAPI + Spark	ROC-AUC 0.845 (from `meta.json`)

Surface	Verification	Outcome
`data/processed/train.parquet`	Read by Spark big-data demo	16,413,084 rows confirmed
`data/processed/test.parquet`	Read by Spark big-data demo	4,105,036 rows confirmed
`data/processed/{product,reviewer}_aggregates/*`	Served by `/aggregates/*`	1.6 M products, 11.5 M reviewers
`reports/ml/drift_report.json`	Surfaced through `/metadata`	Present and well-formed
`reports/ml/threshold_study.csv`	Surfaced through `/metadata`	Present, 9 threshold steps
`reports/ml/fraud_calibration_{raw,sigmoid,isotonic}.png`	Read by `build_report`	Present (raw / sigmoid / isotonic)
`reports/ml/{sentiment,fraud}_error_analysis.csv`	Used by `error_analysis.py`	188 + 85 MB; gitignored
FastAPI `GET /healthz`	curl	**200 OK**
FastAPI `GET /metadata`	curl	Returns full metadata payload incl. selection + drift
FastAPI `POST /predict`	curl + form	Sentiment + fraud_proba + LLM `fraud_explanation`
FastAPI `POST /predict/batch`	curl with 3-row batch	All 3 scored, returns `{predictions: [...]}`
FastAPI `GET /aggregates/products`	curl	Top product `B01415QHYW` (42,644 reviews, 11.96% fraud-rate)
FastAPI `GET /aggregates/fraud-reviewers`	curl	Top suspicious reviewer with 66 reviews, fraud_rate 1.00
FastAPI `GET /stream/recent`	curl	Returns Kafka-sourced scored rows
Spark Structured Streaming (Kafka source)	Producer + consumer end-to-end	Verified 113 → 118 micro-batch files written
Streamlit demo	Boot + `/_stcore/health`	200 OK on `:8501`
MCP server	Boot via stdio	`Amazon Reviews — Sentiment & Fraud Detector` started, 3 tools registered
Ollama review auditor	`make audit PRODUCT=B01415QHYW`	Audit completed; HIGH risk verdict; 4 tool calls (`get_product_aggregate`, `get_product_reviews`, `get_top_fraud_reviewers`, `score_review`)
Ollama LLM explanation	Form submission	`llama3.1:8b` produced realistic risk summaries on real text
Kafka broker (homebrew kafka)	`kafka-get-offsets` + topic describe	Online; topic `reviews` partition 0
Kafka UI (provectus/kafka-ui in Docker)	Browser at `:8090`	Connected to broker; live message viewer working

Every component listed above was either verified or measured in this session, against the same persisted models and data described elsewhere in the report.

18.2 Spark batch on 20.5 M rows – live demo

A new helper, `scripts/spark_bigdata_demo.py`, reads the full **20.5 M-row** processed corpus (train + test parquets, 3.8 GB on disk) and runs four representative aggregations with shuffles and a join. Spark UI is exposed on `http://localhost:4040` while the job runs.

```
SPARK_DRIVER_MEM=8g python3 -m scripts.spark_bigdata_demo
```

Run results (full 20,518,120 rows, single-node `local[*]`):

Stage	Operation	Shuffle?	Time
Stage 1	`avg(star_rating)`, `count(*)`, `avg(fraud_label)`	No	**1.8 s**
Stage 2	`groupBy(reviewer_id)` -> top-10 by review count	Yes (by `reviewer_id`)	**10.7 s**
Stage 3	`groupBy(product_id)` -> top-5 by fraud rate	Yes (by `product_id`)	**2.2 s**
Stage 4	Inner join + filter + count	Yes (large hash-join)	**13.5 s**
	Total		**30.8 s**

Headline numbers from the Spark UI mid-job:

- **Stage 8** (the reviewer_id groupBy): 45 parallel tasks, **603 MB read**, **626 MB shuffle write**.
- Aggregate results: global avg star_rating = 4.008; global fraud_rate = 0.0357; reviewer with highest count = `AG73BVBKU0H22USSFJA5ZWL7AKXA` (410 reviews, fraud_rate 0.122); product with highest fraud rate >= 200 reviews = `B00BI9AKJI` (363 reviews, fraud_rate 0.413).
- 237,967 reviews are on products with fraud_rate >= 0.20.

These exact numbers are reproducible by running the helper. They show the production-scale Spark mechanics (DAG scheduling, partitioned tasks, shuffle write/read, AQE) that scale unchanged to a multi-node cluster via `spark-submit --master ...`.

18.3 Kafka real-time streaming (variable batch sizes)

The Kafka path was driven through three deliberate phases to expose how Spark Structured Streaming pulls data from Kafka in micro-batches.

Phase 1 – slow drip (1 msg / 3 s for 30 s). Producer rate < trigger rate -> Spark micro-batches range 1-6 rows.

Phase 2 – burst (50 msgs in <1 s). Producer dumps a backlog -> Spark's next 5-second trigger gulps **all 50 rows in one micro-batch** (`kafka-batch-00000118.json`).

Phase 3 – idle (30 s, no producer). Spark's trigger continues to fire (~6 times) but produces no output because Kafka has nothing new since the last commit.

Phase	Producer rate	Spark batch sizes seen
Slow drip	~0.3 msg/s	1, 1, 1, 6
Burst	50 msgs in 1 burst	**50 (all in one micro-batch)**
Idle	0 msg/s	0 (no output written)

Sustained-rate test (separate run, before the demo): producer pushed **2,000 messages in 13 s = 153 msg/s**. Spark drained the backlog in a single 2,000-row micro-batch.

This experimentally confirms the two iron rules of Spark Structured Streaming:

1. **Trigger interval is fixed** (5 s in this project, configurable via `--trigger-seconds`). Spark wakes up every interval regardless of work.
2. **Batch size is variable.** Whatever Kafka has accumulated since the last committed offset is what gets pulled.

18.4 Form -> Kafka -> Spark integration

In commit `bb10288`, `POST /predict` and `POST /predict/batch` in `src/serve/app.py` were extended with a fire-and-forget Kafka publish. When the FastAPI process is started with `PUBLISH_PREDICT_TO_KAFKA=1`, every form submission is *also* published to the `reviews` topic, so it flows through Kafka -> Spark Structured Streaming -> `/stream/recent` – in addition to returning the synchronous result inline.

Verified live:

```
BEFORE submit:  Kafka offset=2667
form submit:    POST /predict body="thsi is not so good"
sync response:  sentiment=neutral  fraud_proba=0.344  flag=0
AFTER submit:   Kafka offset=2668  (+1, the form text is now in Kafka)
~5 s later:     Spark micro-batch 112 picks it up, scores it
                kafka-batch-00000112.json contains the same review_body,
                sentiment=neutral, fraud_proba=0.3442
```

The fraud_proba is identical (`0.344` vs `0.3442`) because both paths use the same `models/fraud_pipeline.joblib` – only the transport differs.

Architecturally, this is the fan-out pattern: `/predict` remains a low-latency synchronous API for users, while the same call also feeds the persistent streaming path for downstream consumers (analytics, alerting, archival, etc.).

The producer client is created lazily on first call, with `acks=1, linger_ms=0`. `acks=0` was tried first but messages were silently dropped under this broker's KRaft config; `acks=1` is the correct setting for a one-broker dev setup.

The publish helper is gated by an env var so the patch is safe even when Kafka is not running – `_maybe_get_kafka_producer()` returns `None` if the variable is unset, and `_publish_review_to_kafka()` exits early.

18.5 Three-window observability stack

For the live demo, three browser-based UIs were brought up so the same data can be inspected from three different lenses simultaneously.

URL	UI	What it shows
`http://localhost:8090`	**Kafka UI** (`provectuslabs/kafka-ui` in Docker)	Brokers, topics, partition stats, **live message viewer** with offset / timestamp / JSON value, consumer groups
`http://localhost:4040`	**Spark UI**	Jobs (one per micro-batch), stages, tasks, DAG visualisation, **Structured Streaming tab** with input rate / process rate / batch duration time-series, executor memory
`http://127.0.0.1:8000`	**FastAPI dashboard**	KPI cards from `/metadata`, threshold-study chart, score-a-review form with LLM explanation, **Live Stream tab** that polls `/stream/recent`

Producer: `python3 -m scripts.kafka_producer ...` or simply submitting the form on the FastAPI dashboard.

This is the closest the project gets to a real "control plane": three independent browser tabs, each backed by a different process, watching the same review flow through.

Kafka UI setup (one-time)

`provectuslabs/kafka-ui` runs in Docker and connects to the host's Kafka broker. Because Docker Desktop on macOS isolates the container's network, the broker's default `advertised.listeners=PLAINTEXT://localhost:9092` is unreachable from inside the container (`localhost` resolves to the container itself). The fix is a **dual-listener** Kafka config – add a second listener that advertises a Docker-reachable hostname:

```
# /opt/homebrew/etc/kafka/server.properties
listeners=PLAINTEXT://:9092,CONTROLLER://:9093,DOCKER://:9094
advertised.listeners=PLAINTEXT://localhost:9092,CONTROLLER://localhost:9093,DOCKER://h
listener.security.protocol.map=CONTROLLER:PLAINTEXT,PLAINTEXT:PLAINTEXT,DOCKER:PLAINTE
```

Then run Kafka UI pointed at the DOCKER listener:

```
docker run -d --name kafka-ui --rm -p 8090:8080 \
  -e KAFKA_CLUSTERS_0_NAME=local \
  -e KAFKA_CLUSTERS_0_BOOTSTRAPSERVERS=host.docker.internal:9094 \
  provectuslabs/kafka-ui:latest
```

Existing host clients (Spark consumer, `kafka_producer.py`, `publish_one_review.py`, FastAPI's fire-and-forget publish) continue to use the original `localhost:9092` listener with no code change.

18.6 Bug found, not yet fixed

`src/stream/score_stream.py:74` has an incorrect Spark API call:

```
F.cast(F.col("value"), StringType()).alias("raw_json")  # raises AttributeError
```

The correct form is `F.col("value").cast(StringType()).alias(...)`. The sibling `src/stream/score_stream_kafka.py:79` uses the right call and is what `make stream-kafka` invokes – that is the working path verified end-to-end. The broken `score_stream.py` is reachable via `make stream` and `make stream-once`.

A one-line fix is straightforward and is logged as a follow-up.

18.7 Bugs fixed and code added

`scripts/kafka_producer.py` – streaming line read

Previous behaviour: `fh.readlines()` slurped the entire 11 GB `reviews.jsonl` into Python memory before sending the first message, causing a hang at startup (the producer process accumulated 435 GB of virtual memory before being killed).

After the fix (commit `bb10288`):

```
def line_stream():
    while True:
        with args.source.open() as fh:
            for line in fh:
                yield line
```

Generator yields one line at a time and cycles when exhausted. The producer now scales to any `--n-batches` and `--batch-size`, regardless of source file size. Verified at 100 msgs/burst, 20 bursts, no sleep = **2,000 msgs in 13 s**.

`scripts/kafka_tap.py` (new)

Live tap on the `reviews` topic. Subscribes via `kafka-python` from the latest offset and prints each arriving message in a one-line readable format:

```
[in→kafka offset=2295] product=B09DCSNDT3  rating=5  body="Item came as described! Fast sh
```

Useful as a monitor while running producer + Spark consumer, without needing to start `kafka-console-consumer`.

`scripts/publish_one_review.py` (new)

One-shot helper that publishes a single user-supplied review to the topic (with auto-generated `review_id`, `product_id`, `reviewer_id`). Used to push concrete reviews through the pipeline on demand:

```
python3 -m scripts.publish_one_review \
  --body "Hi, I have used the Chanel perfume and it is awesome" \
  --rating 5 --product BUSER_CHANEL01 --reviewer R_DEMO_RAJESH
```

`scripts/spark_bigdata_demo.py` (new)

The 4-stage Spark aggregation demo described in Section 18.2.

`src/serve/app.py` – fire-and-forget Kafka publish

The form-to-Kafka wiring described in Section 18.4. Gated by `PUBLISH_PREDICT_TO_KAFKA=1`.

18.8 Property snapshot at end of session

Kafka

Property	Value
Broker count	1 (KRaft mode, `node.id=1`)
Listeners	`PLAINTEXT://:9092`, `CONTROLLER://:9093`, `DOCKER://:9094`
Advertised listeners	`localhost:9092` (host clients) + `host.docker.internal:9094` (Docker containers)
Topic	`reviews`
Partitions	1
Replication factor	1
Log retention	168 h (7 days)
Segment max	1 GB
Log dir	`/opt/homebrew/var/lib/kraft-combined-logs/`
Total messages stored at end of session	2,729
Avg msg size	656 bytes
Log file size	1.75 MB

Spark Structured Streaming

Property	Value
App name	`amazon-kafka-scorer`
Trigger interval	5 s
`startingOffsets`	`latest` (only new msgs)
`spark.sql.shuffle.partitions`	4 (streaming) / 200 (batch)
`spark.driver.memory`	8 GB
AQE	enabled
Checkpoint dir	`data/streaming_out/_checkpoints_kafka/`

Producer (form fire-and-forget)

Property	Value
`acks`	1
`linger_ms`	0
`request_timeout_ms`	2000

Measured throughput / latency

Metric	Value
Kafka publish rate (single producer)	**153 msg/s** sustained
Largest single Spark micro-batch	**2,000 rows**
Form -> Kafka -> Spark -> /stream/recent end-to-end	~5 s
20.5 M-row Spark aggregation (4 stages)	**30.8 s**
Heaviest stage: Stage 8 (groupBy reviewer_id)	45 tasks, 603 MB read, 626 MB shuffle write

---

Appendix A – Code changes

File	Change	Commit
`src/common/spark.py`	Driver memory 2g -> 8g (env-overridable), shuffle partitions 8 -> 200, AQE on, Kryo serializer, larger maxResultSize.	`8e2c65c`
`src/train/train.py`	Column-subset parquet read + dtype downcast; sentiment LogReg switched to `saga`; replaced `GradientBoostingClassifier` and heavy `RandomForestClassifier` with `SGDClassifier(log_loss)` and `LogisticRegression(saga)`; tightened TF-IDF (`min_df=10`, fraud `max_features=15_000`).	`10b63b7`
`streamlit_app.py` (new) + `Makefile`	Single-file Streamlit dashboard + `make streamlit` target.	`a9431fc`
`PROJECT_REPORT.md` (new) + `PROJECT_REPORT.pdf`	Build report.	`1e0f013`, `d267887`, then expanded

File	Change	Commit
README.md	Refreshed for full-scale run + LLM/agent surfaces.	9917974
src/ingest/import_amazon_reviews.py	Streaming adapter: McAuley-Lab schema -> project schema.	9917974
architecture.{png,svg} + er_diagram.{png,svg}	Architecture and ER diagrams committed alongside reports.	b99d5a5
reports/ml/{drift_report.json,confusion_matrix.csv,calibration_*}	Diagnostic outputs from the full-scale run.	b99d5a5
models/thresholds.json	Updated to best_f1_threshold=0.80 from the threshold-tuning study.	b99d5a5
scripts/kafka_producer.py	**Bug fix:** streams JSONL line-by-line via generator instead of readlines()-into-RAM (Section 18.7).	bb10288
src/serve/app.py	**Feature:** optional fire-and-forget Kafka publish on /predict and /predict/batch, gated by PUB–LISH_PREDICT_TO_KAFKA=1 (Section 18.4).	bb10288
scripts/kafka_tap.py (new)	Live tap on Kafka topic; prints each arriving message (Section 18.7).	bb10288
scripts/publish_one_review.py (new)	One-shot CLI helper to publish a review to Kafka (Section 18.7).	bb10288
scripts/spark_bigdata_demo.py (new)	20.5 M-row Spark aggregation demo with shuffles + join (Section 18.2).	bb10288
data/raw/reviews.jsonl	Symlink to the converted 11 GB JSONL (avoids duplicating data).	(not in git – local)
/opt/homebrew/etc/kafka/server.properties	**Broker config:** added DOCKER://:9094 listener advertised as host.docker.internal:9094 so Kafka UI in Docker can reach the broker (Section 18.5).	(not in repo – local broker config)

Appendix B – Commits pushed to `origin/main` (this session)

```
bb10288  feat: wire /predict to Kafka and add Spark/Kafka demo helpers
b99d5a5  docs: add architecture/ER diagrams, ML reports, updated thresholds
9917974  docs: README reflects full real-data run + LLM/agent surfaces
d267887  docs: report covers full-feature run (Kafka + diagnostics + agents)
1e0f013  docs: full real-data run report (PROJECT_REPORT.md + .pdf)
a9431fc  feat: Streamlit demo dashboard
10b63b7  refactor(train): scale trainer to 20M-row corpus
8e2c65c  chore(spark): tune SparkSession for full-scale ETL
```

Browse: https://github.com/Paruchuri-Rajesh/Large_Scale_Product-Review/commits/main

Appendix C – File inventory at end of run

Source code (in git):

```
src/ingest/import_amazon_real.py          ← schema adapter
src/ingest/generate_sample.py             ← synthetic generator (dev)
src/etl/batch_etl.py                      ← Spark batch ETL
src/train/train.py                        ← main trainer
src/train/{features,registry,evaluate,utils}.py
src/train/{baselines,threshold_tuning,error_analysis,
           drift_monitor,ablation_study,calibration_report}.py
src/stream/score_stream.py                ← Kafka consumer
src/stream/score_stream_kafka.py          ← Kafka consumer (alt)
src/serve/app.py                          ← FastAPI
src/serve/fraud_explain.py                ← Ollama + fallback
src/serve/mcp_server.py                   ← MCP server
src/serve/templates/index.html
src/serve/static/{app.js,style.css}
src/agents/review_auditor.py              ← Ollama tool-calling agent
src/common/{config,schema,spark,text}.py
scripts/{feed_stream.py,kafka_producer.py,kafka_tap.py,
         publish_one_review.py,spark_bigdata_demo.py,
         run_pipeline.sh,build_report.py}
streamlit_app.py                          ← Streamlit demo
tests/test_pipeline.py
Makefile / requirements.txt / .gitignore
README.md / PROJECT_REPORT.{md,pdf}
architecture.{png,svg} / er_diagram.{png,svg}
```

Generated artifacts (in `.gitignore`):

```
data/raw/reviews.jsonl                    ← 11 GB project-schema JSONL
data/processed/train.parquet              ← 3.0 GB
data/processed/test.parquet               ← 779 MB
data/processed/product_agg.parquet        ← 17 MB
data/processed/reviewer_agg.parquet       ← 325 MB
```

```
data/streaming_out/scored/kafka-batch-*.json
models/sentiment_pipeline.joblib          ← 3 MB
models/fraud_pipeline.joblib              ← 687 KB
models/meta.json
models/thresholds.json                    ← best_f1_threshold = 0.80
mlruns/                                   ← MLflow tracking
reports/ml/drift_report.json
reports/ml/threshold_study.csv
reports/ml/fraud_calibration_{report,bins,summary}.{csv,json}
reports/ml/{sentiment,fraud}_error_samples.csv  (188 + 85 MB)
```

Appendix D – Glossary

- **Adapter** (`import_amazon_real.py`) – script that converts McAuley-Lab JSONL into project schema, streaming line-by-line.
- **AQE** – Adaptive Query Execution; Spark feature that re-plans shuffles based on runtime statistics.
- **ASIN** – Amazon Standard Identification Number; used as `product_id`.
- **Brier score** – mean squared error of probabilistic predictions; lower is better-calibrated.
- **ECE** – Expected Calibration Error; how far predicted probabilities deviate from observed frequencies.
- **ETL** – Extract, Transform, Load – the pattern of reading from a source, transforming it, and persisting to a queryable store.
- **`fastmcp`** – Python MCP framework used by `src/serve/mcp_server.py`.
- **`foreachBatch`** – Spark Structured Streaming sink that hands each micro-batch to a Python function as a DataFrame.
- **Heuristic / weak label** – label generated by rules rather than human ground truth; used here for fraud.
- **Isotonic calibration** – non-parametric monotone mapping of raw probabilities to better-calibrated ones.
- **KRaft** – Kafka's "Raft" mode that removes the Zookeeper dependency.
- **MCP** – Model Context Protocol; lets external apps (e.g. Claude Desktop) call your tools over JSON-RPC stdio.
- **MLflow** – experiment tracking server; stores params, metrics, artifacts.
- **`saga` solver** – Stochastic Average Gradient Augmented; the scikit-learn LR solver designed for large sparse problems.
- **Window function** – Spark SQL operation that computes aggregates over a partition (e.g. all rows with the same `reviewer_id`).